Yatong Bai

Ph.D. Candidate in Robust Deep Learning, Audio Generation, and Optimization

yatong_bai@berkeley.edu
Website: bai-yt.github.io
LinkedIn: linkedin.com/in/yatong-bai

EDUCATION

University of California, Berkeley, Doctor of Philosophy & Master of Science *Aug 2020 – Summer 2025 (Expected)*
- PhD expected May 2025, MS conferred May 2022.   **GPA:** 4.00 / 4.00
- **Areas:** Deep learning, generative models, audio/music AI, adversarial robustness, optimization, and control.
- **Advisor:** Somayeh Sojoudi.
- **Courses:** Convex optimization and algorithms, Deep neural networks, Statistical learning theory, Deep reinforcement learning, Advanced control systems, Theoretical statistics.

Georgia Institute of Technology, Bachelor of Science   *Aug 2016 – Aug 2020*
- **Double major** in Computer Engineering and Mechanical Engineering.   **GPA:** 4.00 / 4.00
- **Courses:** Machine learning, Computer vision, Signals and systems, Embedded systems, Computer architecture.

GRADUATE-LEVEL WORK EXPERIENCE (For Berkeley experiences please see publications)

Adobe Research, Intern Research Scientist   *San Francisco, CA, May 2024 – Nov 2024*
- Improve AI music generation with reinforcement learning via various reward signals including human preference. Paper soon.

Microsoft Applied Science Group, Research Intern   *Redmond, WA, May 2023 – Aug 2023*
- Accelerate diffusion-model-based text-to-audio generation hundreds of times with minimal performance decrease via distillation.
- Innovatively combine the "consistency distillation" framework and "classifier-free guidance" to distill high-performance models.
- The distilled model needs only one neural net query (diffusion models need 400), enabling improving audio semantics by optimizing audio-space losses. We use the CLAP score loss as an example and use objective and subjective evaluation to show its effectiveness.
- Published as "ConsistencyTTA" at arxiv.org/abs/2309.10740. Project website and example generations at consistency-tta.github.io.
- Code available at github.com/Bai-YT/ConsistencyTTA. Model checkpoints at huggingface.co/Bai-YT/ConsistencyTTA.
- Live demo at huggingface.co/spaces/Bai-YT/ConsistencyTTA.

Scale AI, Machine Learning Research Intern   *San Francisco, CA, May 2022 – Dec 2022*
- Researched proposing a dataset with 15 million image-caption pairs and processed its captions with various language models.
- Evaluated the dataset and provided benchmark results with supervised and self-supervised image classification, object detection, image reconstruction, and generation methods (in PyTorch); visualized the embedding space with dimension reduction methods.
- Used these results to characterize the distribution shift of our data from existing datasets. Preprint paper arxiv.org/abs/2401.04575.

PUBLICATIONS AND PREPRINTS

ConsistencyTTA: Accelerating Diffusion-Based Text-to-Audio Generation with Consistency Distillation
Yatong Bai, T Dang, D Tran, K Koishida, and S Sojoudi. In *INTERSPEECH*, 2024. arxiv.org/abs/2309.10740
- See the Microsoft internship (listed above) for details.

MixedNUTS: Training-Free Accuracy-Robustness Balance via Nonlinearly Mixed Classifiers
Yatong Bai, M Zhou, VM Patel, and S Sojoudi. In *Transactions on Machine Learning Research (TMLR)*, 2024. arxiv.org/abs/2402.02263
- Balances neural network classifiers' (un-attacked) clean accuracy and adversarial robustness without additional training, reducing error rate by up to 31% compared to a standalone model.
- Achieved by proposing "nonlinear base model logit transformations" for "mixed classifiers (introduced in the two entries below)".
- The transformations augment the robust base classifier's benign confidence property, thereby balancing accuracy and robustness.

Ranking Manipulation for Conversational Search Engines
S Pfrommer*, Yatong Bai*, T Gautam, and S Sojoudi (*equal contribution). In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024. arxiv.org/abs/2406.03589
- We consider conversational search engines (CSE), powered by large language models (LLM) retrieval augmented generation (RAG).
- When prompted to recommend websites, CSEs simultaneously pay attention to website topic, content, and location in context.
- By embedding spurious characters in HTML code, a website can deceive CSEs to boost its recommendation ranking in the response.
- Proposes RAGDOLL, an e-commerce website dataset for evaluation.
- Dataset available at huggingface.co/datasets/Bai-YT/RAGDOLL. Project code available at github.com/spfrommer/cse-ranking-manipulation.

Improving the Accuracy-Robustness Trade-Off of Classifiers via Local Adaptive Smoothing
Yatong Bai, BG Anderson, A Kim, and S Sojoudi. In *SIAM Journal on Mathematics of Data Science (SIMODS)*, 2024. arxiv.org/abs/2301.12554
- Based on the work below, we further introduce a "mixing network" to adjust the mixing strengths for benign and attacked inputs differently, further improving the accuracy-robustness trade-off. As of submission, on CIFAR-10 and CIFAR-100 datasets, adaptive smoothing is the second most robust method on RobustBench, with noticeably higher clean accuracy than all other works.
- Project code available at github.com/Bai-YT/AdaptiveSmoothing.

Mixing Classifiers to Alleviate the Accuracy-Robustness Trade-Off

Yatong Bai, BG Anderson, and S Sojoudi. In *Learning for Dynamics & Control Conference (L4DC)*, 2024. arxiv.org/abs/2311.15165

- By building a "mixed classifier", specifically mixing the output probabilities of an accurate (often non-robust) classifier and a robust classifier, we greatly alleviate the accuracy-(adversarial) robustness trade-off and achieve certified robustness.
- Our analysis shows that the robust base classifier (RBC)'s prediction confidence is the main source of this improvement. Specifically, the RBC is more confident in correct examples than incorrect ones. The mixed classifier thus puts more trust in correct predictions.

Efficient Global Optimization of Two-Layer ReLU Networks: Adversarial Training and Quadratic-Time Algorithms

Yatong Bai, T Gautam, and S Sojoudi. In *SIAM Journal on Mathematics of Data Science (SIMODS)*, 2022. arxiv.org/abs/2201.01965

- 2021 INFORMS Data Mining Best Paper Competition (Student Track) Runner-up (2nd out of 48 papers).
- We develop efficient ADMM algorithms for the "convex training" formulation, which trains one-hidden-layer neural networks via convex optimization. We prove that the proposed algorithms polynomially improve the computational complexity.

Initial State Interventions for Deconfounded Imitation Learning

S Pfrommer, **Yatong Bai**, H Lee, and S Sojoudi. In *IEEE Conference on Decision and Control (CDC)*, 2023. arxiv.org/abs/2307.15980

- Imitation learning agents suffer from causal confusion. We use a beta-VAE neural net to obtain disentangled latent representations underlying the observations, and use a statistical test to mask confounding latent variables so that the agent performs significantly better when the observations are confounded.

Practical Convex Formulation of Robust One-Hidden-Layer Neural Network Training

Yatong Bai, T Gautam, Y Gai, and S Sojoudi. In *American Control Conference (ACC)*, 2022. arxiv.org/abs/2105.12237

- We leverage the duality theory and robust optimization techniques to develop efficient convex optimization formulations that train robust one-hidden-layer ReLU neural networks via adversarial training.
- Our method demonstrates improved adversarial robustness on common datasets, including CIFAR-10.

Let's Go Shopping (LGS) – Web-Scale Image-Text Dataset for Visual Concept Understanding

Yatong Bai, U Garg, A Shanker, H Zhang, S Parajuli, E Bas, I Filipovic, AN Chu, ED Fomitcheva, E Branson, A Kim, S Sojoudi, K Cho. *Preprint*, 2024. arxiv.org/abs/2401.04575

- See the Scale AI internship (listed above) for details.

UNDERGRADUATE EXPERIENCE

Georgia Institute of Technology
Undergraduate Student Researcher

TINKER Group, RoboMed Group, Meaud Research Group, GT Off-road — *Jan 2018 – Jan 2020*

- Compiled the SPEC 2017 computer architecture benchmark into ARM binary programs using GCC-ARM; Used the Gem5 computer architecture simulator (in C++) to convert the binary programs into debug trace files.
- Built Graphical User Interfaces (GUIs) for a cochlear dynamics simulator in MATLAB. The GUIs controlled simulations, logged and processed experiment data, and visualized the simulation results.

Senior design project: Avionics and test stand controller for a "Monocopter" aircraft

- Implemented the avionics system of a novel unmanned "Monocopter" and a PID-controlled testbed using C++. The avionics filtered noisy magnetometer readings to accurately recover aircraft heading and control the actuators. Also developed a Windows C# GUI.

Honda Aircraft Company, Engineering Intern
Greensboro, NC, May 2019 – Aug 2019

- Performed dynamic simulations for flap linkages in MSC ADAMS.
- Evaluated the stress, deflection, and kinematics in CATIA via Finite Element Analyses (FEA).
- Defined the flap skew and asymmetry warning thresholds and designed a flap control logic in MATLAB.

Tesla, Inc., Engineering Intern
Palo Alto, CA, May 2018 – Aug 2018

- Implement scripts that convert simulation models between different tolerance stack-up (GD&T) simulators.

ACADEMIC ACTIVITIES

- **Reviewing:** CDC (2022, 2023), ICML (2023, 2024), CCTA (2023), NeurIPS (2023, 2024), and ICLR (2024, 2025) conferences.
- **Teaching:** Graduate Student Instructor (TA) for Spring 22, Fall 22, Fall 23 "IEOR 160: Nonlinear and Discrete Optimization". Graduate Student Instructor (TA) for Fall 24 "EECS 127: Optimization Models in Engineering".
- **Talks/Posters:** INTERSPEECH 2024, L4DC 2024, ACC 2022, CCTA 2023, INFORMS 2021, and MOPTA 2021 conferences.

AWARDS

INFORMS Data Mining Best Paper Competition (student track) Runner-up Prize	*Oct 2021*
The Henry Lurie Family Fellowship	*May 2024*
UC Berkeley Graduate Division Block Grant Fellowship	*April 2021*
Georgia Tech School of ECE Roger P. Webb ECE Senior Scholar Awards	*April 2021*

CODING

I regularly code in Python (PyTorch) and MATLAB and write in LaTeX, often working on remote cloud computing machines.
Other languages I've used include HTML, C, C++, Java, and R.